

timeseries-forecasting-1

October 11, 2024

1 Climate Data Time-Series

You are again moving to another role, not at *The Weather Channel*, where you are asked to create a Weather Forecasting Model.

For that, you will be using *Jena Climate* dataset recorded by the *Max Planck Institute for Biogeochemistry*.

The dataset consists of 14 features such as temperature, pressure, humidity etc, recorded **once per 10 minutes**.

Location: Weather Station, Max Planck Institute for Biogeochemistry in Jena, Germany

Time-frame Considered: **Jan 10, 2009 - December 31, 2012**

Library Imports

```
[1]: import pandas as pd
import matplotlib.pyplot as plt
import keras
```

1.0.1 1) Load your data

Your data can be found on the Deep Learning Module under a file named: `climate_data_2009_2012.csv`

```
[2]: df = pd.read_csv("climate_data_2009_2012.csv")
```

1.0.2 2) Data engineering

You are given 3 lists: - `titles`: Display names of your columns - `feature_keys`: Names of the columns used as features - `colors`: The color to use when plotting that column's value

```
[3]: titles = [
    "Pressure",
    "Temperature",
    "Temperature in Kelvin",
    "Temperature (dew point)",
    "Relative Humidity",
    "Saturation vapor pressure",
    "Vapor pressure",
```

```

    "Vapor pressure deficit",
    "Specific humidity",
    "Water vapor concentration",
    "Airtight",
    "Wind speed",
    "Maximum wind speed",
    "Wind direction in degrees",
]

feature_keys = [
    "p (mbar)",
    "T (degC)",
    "Tpot (K)",
    "Tdew (degC)",
    "rh (%)",
    "VPmax (mbar)",
    "VPact (mbar)",
    "VPdef (mbar)",
    "sh (g/kg)",
    "H2OC (mmol/mol)",
    "rho (g/m**3)",
    "wv (m/s)",
    "max. wv (m/s)",
    "wd (deg)",
]

colors = [
    "blue",
    "orange",
    "green",
    "red",
    "purple",
    "brown",
    "pink",
    "gray",
    "olive",
    "cyan",
]

```

Let's look at the climate data:

```
[4]: df.head()
```

```
[4]:
```

	Date Time	p (mbar)	T (degC)	Tpot (K)	Tdew (degC)	rh (%)	\
0	01.01.2009 00:10:00	996.52	-8.02	265.40	-8.90	93.3	
1	01.01.2009 00:20:00	996.57	-8.41	265.01	-9.28	93.4	
2	01.01.2009 00:30:00	996.53	-8.51	264.91	-9.31	93.9	

3	01.01.2009 00:40:00	996.51	-8.31	265.12	-9.07	94.2
4	01.01.2009 00:50:00	996.51	-8.27	265.15	-9.04	94.1

	VPmax (mbar)	VPact (mbar)	VPdef (mbar)	sh (g/kg)	H2OC (mmol/mol)	\
0	3.33	3.11	0.22	1.94	3.12	
1	3.23	3.02	0.21	1.89	3.03	
2	3.21	3.01	0.20	1.88	3.02	
3	3.26	3.07	0.19	1.92	3.08	
4	3.27	3.08	0.19	1.92	3.09	

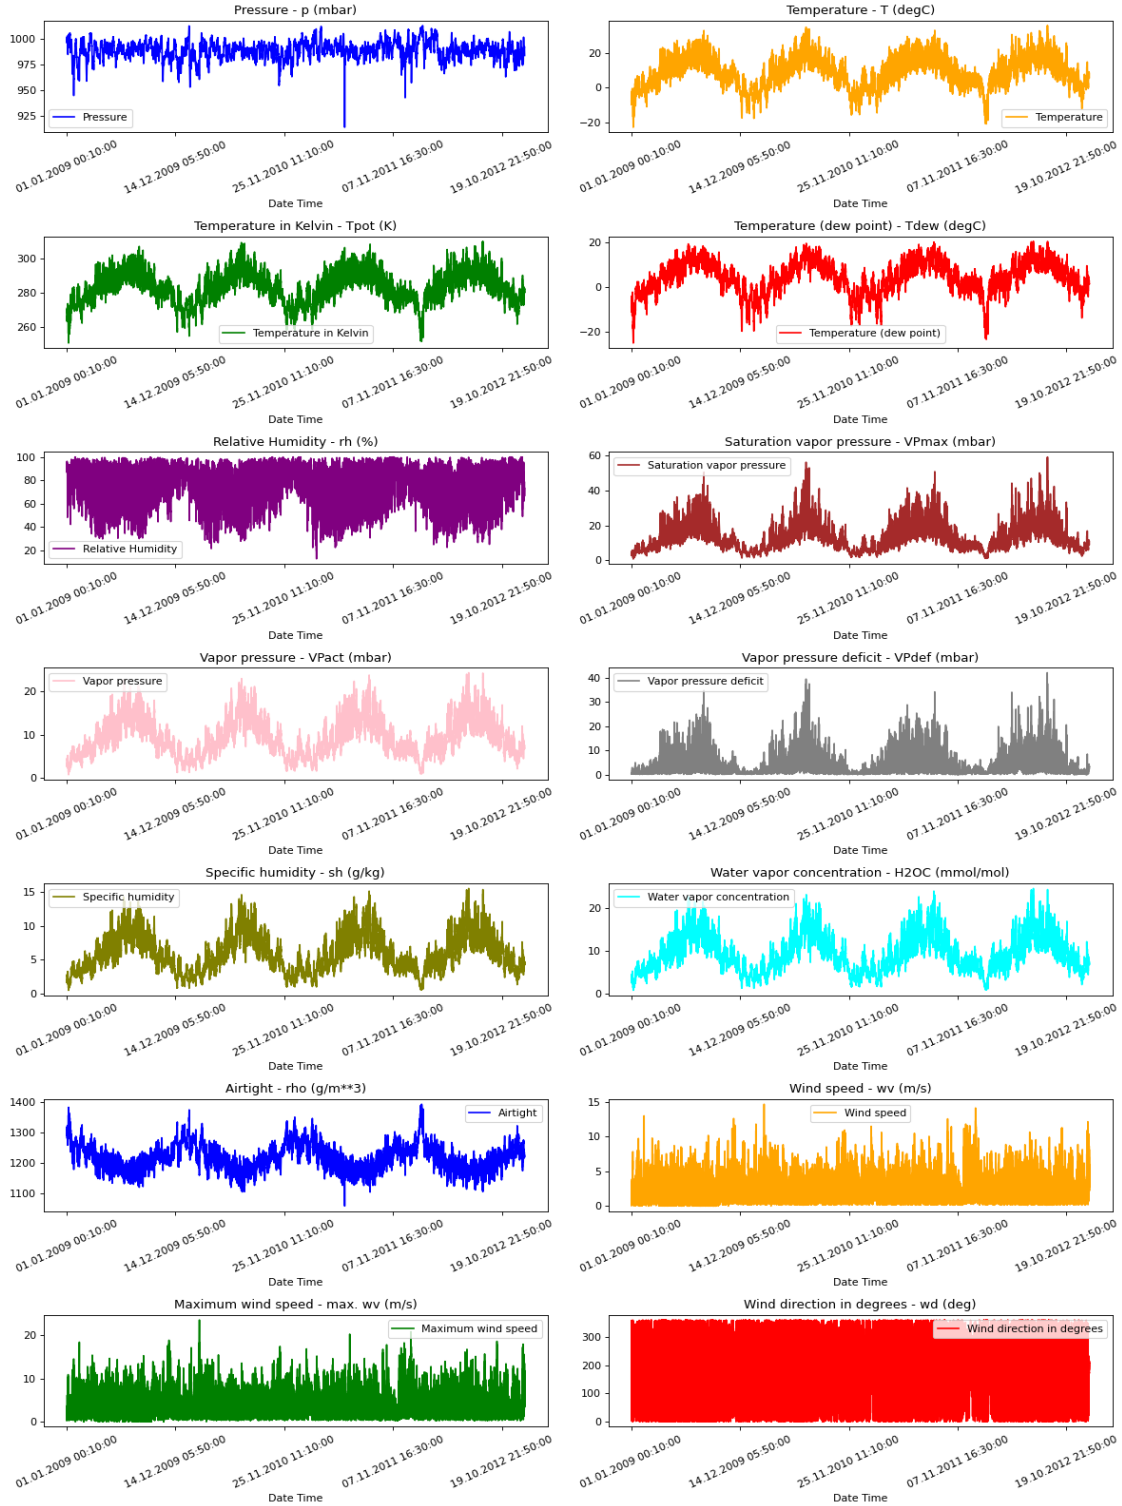
	rho (g/m**3)	wv (m/s)	max. wv (m/s)	wd (deg)
0	1307.75	1.03	1.75	152.3
1	1309.80	0.72	1.50	136.1
2	1310.24	0.19	0.63	171.6
3	1309.19	0.34	0.50	198.0
4	1309.00	0.32	0.63	214.3

Define a function to show a plot of each column (using the respective color)

```
[5]: def show_raw_visualization(data, date_time_key):
    time_data = data[date_time_key]
    fig, axes = plt.subplots(
        nrows=7, ncols=2, figsize=(15, 20), dpi=80, facecolor="w", edgecolor="k"
    )
    for i in range(len(feature_keys)):
        key = feature_keys[i]
        c = colors[i % (len(colors))]
        t_data = data[key]
        t_data.index = time_data
        t_data.head()
        ax = t_data.plot(
            ax=axes[i // 2, i % 2],
            color=c,
            title="{0} - {1}".format(titles[i], key),
            rot=25,
        )
        ax.legend([titles[i]])
    plt.tight_layout()
```

Display each column in a plot using above function:

```
[6]: show_raw_visualization(df, "Date Time")
```



As you can see we have lots of data, this can be a challenge when we train our model, to resolve that we will reduce the resolution of our data, instead of having a climate signal each 10 minutes,

we will have it each hour

- Add a new column to your dataframe with the Date Time information
- Name that column FormatedDateTime
- Convert that column into date time data type
- Set that column as the dataframe index
- Regroup data to be each 1 hour instead of each 10 minutes
- Save the grouped data into a dataframe called df_resampled
- Remove the FormatedDateTime as the index.
- Show the top 5 rows of df_resampled

```
[7]: df['FormatedDateTime'] = pd.to_datetime(df['Date Time'], format='%d.%m.%Y %H:%M:
      ↪%S')
df = df.set_index('FormatedDateTime')
df_resampled = df[feature_keys].resample('H').mean()
df_resampled = df_resampled.reset_index()

df_resampled.head()
```

<ipython-input-7-e61db8939877>:3: FutureWarning: 'H' is deprecated and will be removed in a future version, please use 'h' instead.

```
df_resampled = df[feature_keys].resample('H').mean()
```

```
[7]:
```

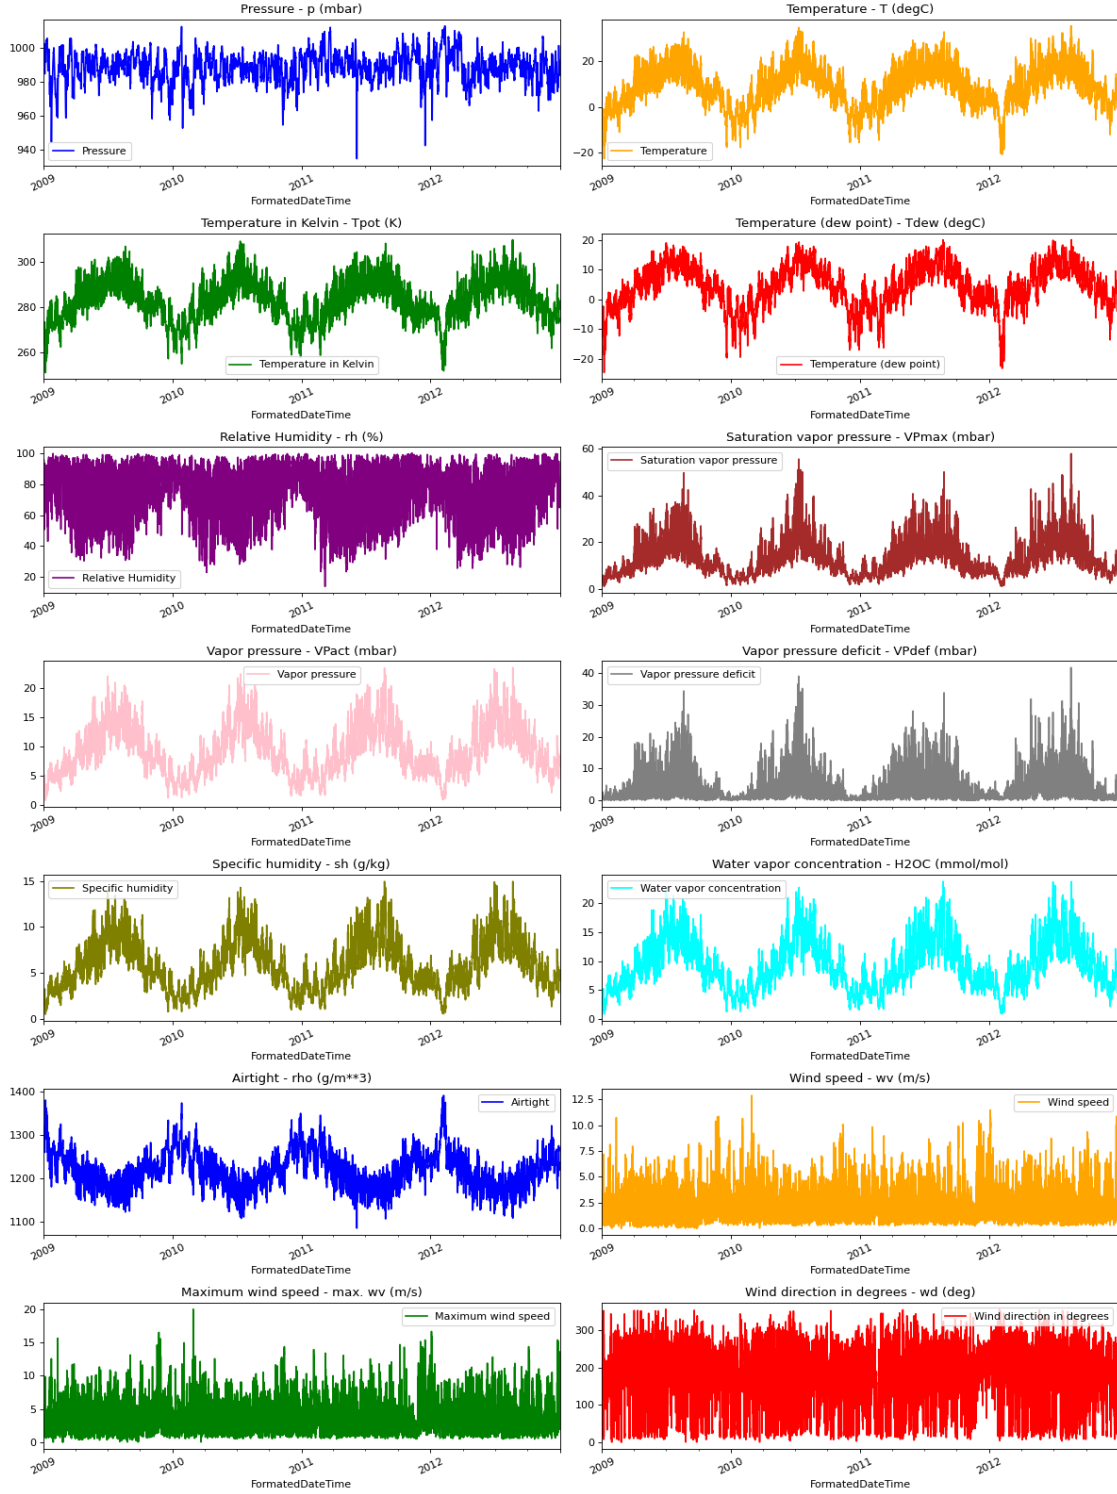
	FormatedDateTime	p (mbar)	T (degC)	Tpot (K)	Tdew (degC)	\
0	2009-01-01 00:00:00	996.528000	-8.304000	265.118000	-9.120000	
1	2009-01-01 01:00:00	996.525000	-8.065000	265.361667	-8.861667	
2	2009-01-01 02:00:00	996.745000	-8.763333	264.645000	-9.610000	
3	2009-01-01 03:00:00	996.986667	-8.896667	264.491667	-9.786667	
4	2009-01-01 04:00:00	997.158333	-9.348333	264.026667	-10.345000	

	rh (%)	VPmax (mbar)	VPact (mbar)	VPdef (mbar)	sh (g/kg)	\
0	93.780000	3.260000	3.058000	0.202000	1.910000	
1	93.933333	3.323333	3.121667	0.201667	1.951667	
2	93.533333	3.145000	2.940000	0.201667	1.836667	
3	93.200000	3.111667	2.898333	0.210000	1.811667	
4	92.383333	3.001667	2.775000	0.231667	1.733333	

	H2OC (mmol/mol)	rho (g/m**3)	wv (m/s)	max. wv (m/s)	wd (deg)
0	3.068000	1309.196000	0.520000	1.002000	174.460000
1	3.133333	1307.981667	0.316667	0.711667	172.416667
2	2.950000	1311.816667	0.248333	0.606667	196.816667
3	2.906667	1312.813333	0.176667	0.606667	157.083333
4	2.780000	1315.355000	0.290000	0.670000	150.093333

Let's look at our fields again

```
[8]: show_raw_visualization(df_resampled, "FormatedDateTime")
```



1.0.3 3) Data Split: Train and Evaluation datasets.

- We are tracking data from past 120 timestamps (120 hours = 5 days).

- This data will be used to predict the temperature after 12 timestamps (12 hours).
- Since every feature has values with varying ranges, we do normalization to confine feature values to a range of [0, 1] before training a neural network.
- We do this by subtracting the mean and dividing by the standard deviation of each feature in the *normalize* function
- The model is shown data for first 5 days i.e. 120 observations, that are sampled every hour.
- The temperature after 12 hours observation will be used as a label.

```
[9]: # 70% of the data will be used for training, the rest for testing
split_fraction = 0.7
# The number of samples is the number of rows in the data
number_of_samples = df_resampled.shape[0]
# The size in rows of the split dataset
train_split = int(split_fraction * int(number_of_samples))

# Number of samples in the past used to predict the future
past = 120
# Number of samples in the future to predict (the value in the 72nd hour is our
↪ label)
future = 12
# Learning rate parameter for the Adam optimizer
learning_rate = 0.001
# Batch size for the model training
batch_size = 256
# Number of epochs for the model training
epochs = 10

# Another way to normalize the data (all columns in the same range)
def normalize(data, train_split):
    data_mean = data[:train_split].mean(axis=0)
    data_std = data[:train_split].std(axis=0)
    return (data - data_mean) / data_std
```

- Let's select the following parameters as our features:
 - Pressure, Temperature, Saturation vapor pressure, Vapor pressure deficit, Specific humidity, Airtight, Wind speed
- Set the column FormatedDateTime as the index of our dataframe.
 - This is important since now, FormatedDateTime is used as our datetime field and not as a Feature field
- Normalize all fields
- Generate two datasets:
 - train_data: Train dataset with our normalized fields
 - val_data: Validation dataset

```
[10]: print(
    "The selected parameters are:",
    ", ".join([titles[i] for i in [0, 1, 5, 7, 8, 10, 11]]),
)
selected_features = [feature_keys[i] for i in [0, 1, 5, 7, 8, 10, 11]]
features = df_resampled[selected_features]
features.index = df_resampled["FormattedDateTime"]
print(features.head())

features = normalize(features.values, train_split)
features = pd.DataFrame(features)
print(features.head())

train_data = features.loc[0 : train_split - 1]
val_data = features.loc[train_split:]
```

The selected parameters are: Pressure, Temperature, Saturation vapor pressure, Vapor pressure deficit, Specific humidity, Airtight, Wind speed

	p (mbar)	T (degC)	VPmax (mbar)	VPdef (mbar)	\
FormattedDateTime					
2009-01-01 00:00:00	996.528000	-8.304000	3.260000	0.202000	
2009-01-01 01:00:00	996.525000	-8.065000	3.323333	0.201667	
2009-01-01 02:00:00	996.745000	-8.763333	3.145000	0.201667	
2009-01-01 03:00:00	996.986667	-8.896667	3.111667	0.210000	
2009-01-01 04:00:00	997.158333	-9.348333	3.001667	0.231667	

	sh (g/kg)	rho (g/m**3)	wv (m/s)
FormattedDateTime			
2009-01-01 00:00:00	1.910000	1309.196000	0.520000
2009-01-01 01:00:00	1.951667	1307.981667	0.316667
2009-01-01 02:00:00	1.836667	1311.816667	0.248333
2009-01-01 03:00:00	1.811667	1312.813333	0.176667
2009-01-01 04:00:00	1.733333	1315.355000	0.290000

	0	1	2	3	4	5	6
0	0.988366	-1.936957	-1.314750	-0.797292	-1.472751	2.198783	-1.116409
1	0.988002	-1.909978	-1.306369	-0.797363	-1.457136	2.169559	-1.256715
2	1.014643	-1.988807	-1.329968	-0.797363	-1.500234	2.261854	-1.303867
3	1.043907	-2.003858	-1.334379	-0.795594	-1.509604	2.285840	-1.353320
4	1.064694	-2.054843	-1.348935	-0.790994	-1.538961	2.347009	-1.275116

Now, here we need to set our Label Dataset.

- We want to use the last 5 days of data, to predict the next 12 hours
- This means that our label starts at the 12th hour after the history data.
 - [..... .]
 - —Start—>
- And it will end at the end of our train dataset size.
 - <— Train —> <— Test —>


```

- [.....|.....]
- -----End----->

```

```

[11]: start = past + future
      end = start + train_split

      x_train = train_data[[i for i in range(7)]].values
      y_train = features.iloc[start:end][[1]]

      step = 1
      sequence_length = past

```

The *timeseries_dataset_from_array* function takes in a sequence of data-points gathered at equal intervals, along with time series parameters such as length of the sequences/windows, spacing between two sequence/windows, etc., to produce batches of sub-timeseries inputs and targets sampled from the main timeseries.

- Input data (hour features) = `x_train`
- The **corresponding** value of the temperature 12 hours into the future = `y_train`
- Since we want to use 5 days of data to predict the future temperature then: `sequence_length = 120`
- Since we want to sample every hour then: `sampling_rate = 1`
- Let's use a common batch size of 256 (variable above)

```

[12]: dataset_train = keras.preprocessing.timeseries_dataset_from_array(
      x_train,
      y_train,
      sequence_length=sequence_length,
      sampling_rate=step,
      batch_size=batch_size,
      )

```

Now let's prepare our validation dataset:

- The validation dataset must not contain the last 120+12 rows as we won't have label data for those records, hence these rows must be subtracted from the end of the data.
- The validation label dataset must start from 120+12 after `train_split`, hence we must add `past + future` to `label_start`.

```

[13]: x_end = len(val_data) - past - future

      label_start = train_split + past + future

      x_val = val_data.iloc[:x_end][[i for i in range(7)]].values
      y_val = features.iloc[label_start:][[1]]

      dataset_val = keras.preprocessing.timeseries_dataset_from_array(
          x_val,
          y_val,

```

```

        sequence_length=sequence_length,
        sampling_rate=step,
        batch_size=batch_size,
    )

    for batch in dataset_train.take(1):
        inputs, targets = batch

    print("Input shape:", inputs.numpy().shape)
    print("Target shape:", targets.numpy().shape)

```

Input shape: (256, 120, 7)

Target shape: (256, 1)

1.0.4 4) Define and Compile your model:

- An input layer
- A Long Short-Term Memory Hidden Layer with 32 units. LSTM is a type of recurrent neural network layer that is well-suited for time series data.
- An output Dense Layer (Linear Activation function)

```

[14]: inputs = keras.layers.Input(shape=(inputs.shape[1], inputs.shape[2]))
      lstm_out = keras.layers.LSTM(32)(inputs)
      outputs = keras.layers.Dense(1)(lstm_out)

      model = keras.Model(inputs=inputs, outputs=outputs)
      model.compile(optimizer=keras.optimizers.Adam(learning_rate=learning_rate),
                    ↪loss="mse")
      model.summary()

```

Model: "functional"

Layer (type)	Output Shape	
↪Param #		
input_layer (InputLayer)	(None, 120, 7)	↪
↪ 0		
lstm (LSTM)	(None, 32)	↪
↪5,120		
dense (Dense)	(None, 1)	↪
↪ 33		

Total params: 5,153 (20.13 KB)

Trainable params: 5,153 (20.13 KB)

Non-trainable params: 0 (0.00 B)

1.0.5 5) Train your model:

Specify the file path where the model's weights will be saved with: `path_checkpoint = "model_checkpoint.weights.h5"`

We want to add a callback to stop training when a monitored metric stops improving: `es_callback = keras.callbacks.EarlyStopping(monitor="val_loss", min_delta=0, patience=5)`

Train the model using Fit

```
[15]: path_checkpoint = "model_checkpoint.weights.h5"
es_callback = keras.callbacks.EarlyStopping(monitor="val_loss", min_delta=0,
patience=5)

modelckpt_callback = keras.callbacks.ModelCheckpoint(
    monitor="val_loss",
    filepath=path_checkpoint,
    verbose=1,
    save_weights_only=True,
    save_best_only=True,
)

history = model.fit(
    dataset_train,
    epochs=epochs,
    validation_data=dataset_val,
    callbacks=[es_callback, modelckpt_callback],
)
```

Epoch 1/10

96/96 0s 109ms/step -

loss: 0.7310

Epoch 1: val_loss improved from inf to 0.25503, saving model to
model_checkpoint.weights.h5

96/96 18s 144ms/step -

loss: 0.7278 - val_loss: 0.2550

Epoch 2/10

96/96 0s 114ms/step -

loss: 0.2128

Epoch 2: val_loss improved from 0.25503 to 0.21500, saving model to
model_checkpoint.weights.h5

```

96/96          14s 145ms/step -
loss: 0.2127 - val_loss: 0.2150
Epoch 3/10
96/96          0s 116ms/step -
loss: 0.1676
Epoch 3: val_loss improved from 0.21500 to 0.19015, saving model to
model_checkpoint.weights.h5
96/96          14s 147ms/step -
loss: 0.1676 - val_loss: 0.1902
Epoch 4/10
96/96          0s 115ms/step -
loss: 0.1460
Epoch 4: val_loss improved from 0.19015 to 0.16143, saving model to
model_checkpoint.weights.h5
96/96          14s 146ms/step -
loss: 0.1460 - val_loss: 0.1614
Epoch 5/10
96/96          0s 113ms/step -
loss: 0.1318
Epoch 5: val_loss improved from 0.16143 to 0.14703, saving model to
model_checkpoint.weights.h5
96/96          14s 145ms/step -
loss: 0.1317 - val_loss: 0.1470
Epoch 6/10
95/96          0s 121ms/step -
loss: 0.1250
Epoch 6: val_loss improved from 0.14703 to 0.13824, saving model to
model_checkpoint.weights.h5
96/96          21s 146ms/step -
loss: 0.1250 - val_loss: 0.1382
Epoch 7/10
95/96          0s 122ms/step -
loss: 0.1199
Epoch 7: val_loss improved from 0.13824 to 0.13152, saving model to
model_checkpoint.weights.h5
96/96          14s 147ms/step -
loss: 0.1198 - val_loss: 0.1315
Epoch 8/10
95/96          0s 118ms/step -
loss: 0.1155
Epoch 8: val_loss improved from 0.13152 to 0.12676, saving model to
model_checkpoint.weights.h5
96/96          14s 143ms/step -
loss: 0.1155 - val_loss: 0.1268
Epoch 9/10
95/96          0s 122ms/step -
loss: 0.1124
Epoch 9: val_loss improved from 0.12676 to 0.12350, saving model to


```

```
model_checkpoint.weights.h5
96/96          14s 144ms/step -
loss: 0.1124 - val_loss: 0.1235
Epoch 10/10
96/96          0s 122ms/step -
loss: 0.1103
Epoch 10: val_loss improved from 0.12350 to 0.12106, saving model to
model_checkpoint.weights.h5
96/96          14s 148ms/step -
loss: 0.1103 - val_loss: 0.1211
```

Plot the results of your training:

```
[16]: def visualize_loss(history, title):
    loss = history.history["loss"]
    val_loss = history.history["val_loss"]
    epochs = range(len(loss))
    plt.figure()
    plt.plot(epochs, loss, "b", label="Training loss")
    plt.plot(epochs, val_loss, "r", label="Validation loss")
    plt.title(title)
    plt.xlabel("Epochs")
    plt.ylabel("Loss")
    plt.legend()
    plt.show()

visualize_loss(history, "Training and Validation Loss")
```



Make 5 predictions and display the predicted value

```
[17]: def show_plot(plot_data, delta, title):
    labels = ["History", "True Future", "Model Prediction"]
    marker = [".-", "rx", "go"]
    time_steps = list(range(-(plot_data[0].shape[0]), 0))
    if delta:
        future = delta
    else:
        future = 0

    plt.title(title)
    for i, val in enumerate(plot_data):
        if i:
            plt.plot(future, plot_data[i], marker[i], markersize=10,
↪ label=labels[i])
        else:
            plt.plot(time_steps, plot_data[i].flatten(), marker[i],
↪ label=labels[i])
    plt.legend()
```

```


plt.xlim([time_steps[0], (future + 5) * 2])
plt.xlabel("Time-Step")
plt.show()
return

for x, y in dataset_val.take(5):
    show_plot(
        [x[0][:, 1].numpy(), y[0].numpy(), model.predict(x)[0]],
        12,
        "Single Step Prediction",
    )

```


8/8

0s 12ms/step




8/8

0s 11ms/step




8/8

0s 12ms/step




8/8

0s 12ms/step



8/8

0s 12ms/step



Now make a Time Series Forecasting where using the last 3 days you will predict the weather in the next 3 hours.

```
[18]: split_fraction = 0.7

print(
    "Los parámetros seleccionados son:",
    ", ".join([titles[i] for i in [0, 1, 5, 7, 8, 10, 11]]),
)

selected_features = [feature_keys[i] for i in [0, 1, 5, 7, 8, 10, 11]]
features = df_resampled[selected_features]
features.index = df_resampled["FormattedDateTime"]
print(features.head())

features = normalize(features.values, train_split)
features = pd.DataFrame(features)
print(features.head())

train_data = features.loc[0 : train_split - 1]
```

```

val_data = features.loc[train_split:]

start = past + future
end = start + train_split

x_train = train_data[[i for i in range(7)]].values
y_train = features.iloc[start:end][[1]]

step = 1
sequence_length = past

dataset_train = keras.preprocessing.timeseries_dataset_from_array(
    x_train,
    y_train,
    sequence_length=sequence_length,
    sampling_rate=step,
    batch_size=batch_size,
)

x_end = len(val_data) - past - future
label_start = train_split + past + future

x_val = val_data.iloc[:x_end][[i for i in range(7)]].values
y_val = features.iloc[label_start:][[1]]

dataset_val = keras.preprocessing.timeseries_dataset_from_array(
    x_val,
    y_val,
    sequence_length=sequence_length,
    sampling_rate=step,
    batch_size=batch_size,
)

for batch in dataset_train.take(1):
    inputs, targets = batch

print("Forma de entrada:", inputs.numpy().shape)
print("Forma de salida:", targets.numpy().shape)

inputs = keras.layers.Input(shape=(inputs.shape[1], inputs.shape[2]))
lstm_out = keras.layers.LSTM(32)(inputs)
outputs = keras.layers.Dense(1)(lstm_out)

model = keras.Model(inputs=inputs, outputs=outputs)
model.compile(optimizer=keras.optimizers.Adam(learning_rate=learning_rate),
    ↪ loss="mse")
model.summary()

```

```

path_checkpoint = "model_checkpoint.weights.h5"
es_callback = keras.callbacks.EarlyStopping(monitor="val_loss", min_delta=0,
patience=5)

modelckpt_callback = keras.callbacks.ModelCheckpoint(
    monitor="val_loss",
    filepath=path_checkpoint,
    verbose=1,
    save_weights_only=True,
    save_best_only=True,
)

history = model.fit(
    dataset_train,
    epochs=epochs,
    validation_data=dataset_val,
    callbacks=[es_callback, modelckpt_callback],
)

number_of_samples = df_resampled.shape[0]
train_split = int(split_fraction * int(number_of_samples))

past = 72
future = 3

learning_rate = 0.001
batch_size = 256
epochs = 10

def normalize(data, train_split):
    data_mean = data[:train_split].mean(axis=0)
    data_std = data[:train_split].std(axis=0)
    return (data - data_mean) / data_std

```

Los parámetros seleccionados son: Pressure, Temperature, Saturation vapor pressure, Vapor pressure deficit, Specific humidity, Airtight, Wind speed

	p (mbar)	T (degC)	VPmax (mbar)	VPdef (mbar)	\
FormattedDateTime					
2009-01-01 00:00:00	996.528000	-8.304000	3.260000	0.202000	
2009-01-01 01:00:00	996.525000	-8.065000	3.323333	0.201667	
2009-01-01 02:00:00	996.745000	-8.763333	3.145000	0.201667	
2009-01-01 03:00:00	996.986667	-8.896667	3.111667	0.210000	
2009-01-01 04:00:00	997.158333	-9.348333	3.001667	0.231667	
	sh (g/kg)	rho (g/m**3)	wv (m/s)		
FormattedDateTime					

```

2009-01-01 00:00:00    1.910000    1309.196000    0.520000
2009-01-01 01:00:00    1.951667    1307.981667    0.316667
2009-01-01 02:00:00    1.836667    1311.816667    0.248333
2009-01-01 03:00:00    1.811667    1312.813333    0.176667
2009-01-01 04:00:00    1.733333    1315.355000    0.290000
      0          1          2          3          4          5          6
0  0.988366 -1.936957 -1.314750 -0.797292 -1.472751  2.198783 -1.116409
1  0.988002 -1.909978 -1.306369 -0.797363 -1.457136  2.169559 -1.256715
2  1.014643 -1.988807 -1.329968 -0.797363 -1.500234  2.261854 -1.303867
3  1.043907 -2.003858 -1.334379 -0.795594 -1.509604  2.285840 -1.353320
4  1.064694 -2.054843 -1.348935 -0.790994 -1.538961  2.347009 -1.275116
Forma de entrada: (256, 120, 7)
Forma de salida: (256, 1)
Model: "functional_1"

```

Layer (type)	Output Shape	
Param #		
input_layer_1 (InputLayer)	(None, 120, 7)	
↪ 0		
lstm_1 (LSTM)	(None, 32)	
↪ 5,120		
dense_1 (Dense)	(None, 1)	
↪ 33		

Total params: 5,153 (20.13 KB)

Trainable params: 5,153 (20.13 KB)

Non-trainable params: 0 (0.00 B)

```

Epoch 1/10
95/96          0s 131ms/step -
loss: 0.7063
Epoch 1: val_loss improved from inf to 0.25311, saving model to
model_checkpoint.weights.h5
96/96          18s 158ms/step -
loss: 0.7003 - val_loss: 0.2531
Epoch 2/10
95/96          0s 123ms/step -

```

```

loss: 0.2131
Epoch 2: val_loss improved from 0.25311 to 0.18389, saving model to
model_checkpoint.weights.h5
96/96          14s 148ms/step -
loss: 0.2128 - val_loss: 0.1839
Epoch 3/10
95/96          0s 126ms/step -
loss: 0.1556
Epoch 3: val_loss improved from 0.18389 to 0.15997, saving model to
model_checkpoint.weights.h5
96/96          15s 151ms/step -
loss: 0.1555 - val_loss: 0.1600
Epoch 4/10
95/96          0s 120ms/step -
loss: 0.1399
Epoch 4: val_loss improved from 0.15997 to 0.14783, saving model to
model_checkpoint.weights.h5
96/96          14s 146ms/step -
loss: 0.1398 - val_loss: 0.1478
Epoch 5/10
95/96          0s 121ms/step -
loss: 0.1314
Epoch 5: val_loss improved from 0.14783 to 0.14032, saving model to
model_checkpoint.weights.h5
96/96          14s 145ms/step -
loss: 0.1313 - val_loss: 0.1403
Epoch 6/10
95/96          0s 117ms/step -
loss: 0.1262
Epoch 6: val_loss improved from 0.14032 to 0.13708, saving model to
model_checkpoint.weights.h5
96/96          13s 139ms/step -
loss: 0.1261 - val_loss: 0.1371
Epoch 7/10
95/96          0s 118ms/step -
loss: 0.1235
Epoch 7: val_loss improved from 0.13708 to 0.13408, saving model to
model_checkpoint.weights.h5
96/96          14s 142ms/step -
loss: 0.1234 - val_loss: 0.1341
Epoch 8/10
95/96          0s 119ms/step -
loss: 0.1205
Epoch 8: val_loss improved from 0.13408 to 0.13030, saving model to
model_checkpoint.weights.h5
96/96          14s 142ms/step -
loss: 0.1204 - val_loss: 0.1303
Epoch 9/10

```

```

96/96          0s 131ms/step -
loss: 0.1173
Epoch 9: val_loss improved from 0.13030 to 0.12622, saving model to
model_checkpoint.weights.h5
96/96          15s 158ms/step -
loss: 0.1172 - val_loss: 0.1262
Epoch 10/10
95/96          0s 119ms/step -
loss: 0.1143
Epoch 10: val_loss improved from 0.12622 to 0.12243, saving model to
model_checkpoint.weights.h5
96/96          14s 143ms/step -
loss: 0.1142 - val_loss: 0.1224

```

```

[19]: def visualize_loss(history, title):
    loss = history.history["loss"]
    val_loss = history.history["val_loss"]
    epochs = range(len(loss))

    plt.figure()
    plt.plot(epochs, loss, "b", label="Training loss")
    plt.plot(epochs, val_loss, "r", label="Validation loss")
    plt.title(title)
    plt.xlabel("Epochs")
    plt.ylabel("Loss")
    plt.legend()
    plt.show()

def show_plot(plot_data, delta, title):
    labels = ["History", "True Future", "Model Prediction"]
    marker = [".-", "rx", "go"]
    time_steps = list(range(-(plot_data[0].shape[0]), 0))

    future = delta if delta else 0

    plt.title(title)
    for i, val in enumerate(plot_data):
        if i:
            plt.plot(future, plot_data[i], marker[i], markersize=10,
↪label=labels[i])
        else:
            plt.plot(time_steps, plot_data[i].flatten(), marker[i],
↪label=labels[i])

    plt.legend()
    plt.xlim([time_steps[0], (future + 5) * 2])
    plt.xlabel("Time-Step")

```



```

plt.show()
return


for x, y in dataset_val.take(5):
    show_plot(
        [x[0][:, 1].numpy(), y[0].numpy(), model.predict(x)[0]],
        3,
        "Single Step Prediction",
    )

visualize_loss(history, "Training and Validation Loss")

```


8/8

0s 12ms/step




8/8

0s 12ms/step




8/8

0s 14ms/step



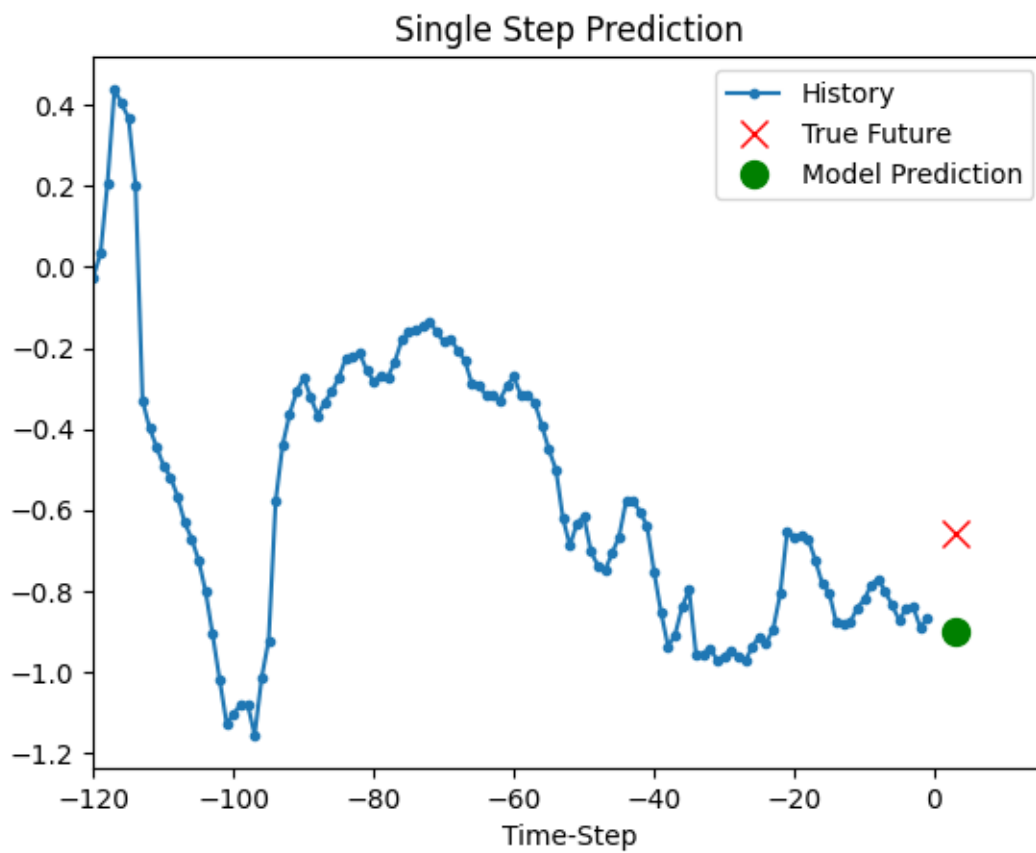
8/8


0s 16ms/step



8/8

0s 19ms/step





Como podemos analizar con los resultados del nuevo modelo creado este funciona razonablemente bien en la predicción que queremos hacer de las próximas 3 horas usando los datos de los últimos 3 días. Podemos ver en las graficas que el modelo captura las tendencias generales El entrenamiento fue efectivo, ya que vemos como las pérdidas de entrenamiento y validación están disminuyendo de manera adecuada.

Los cambios principales en el nuevo modelo para predecir las siguientes 3 horas con los ultimos 3 días fueron:

Se ajustó la duración de los datos historicos, se cambió past de 120 o sea 5 días, a 72 o sea 3 días. Tambien se cambió a futuro, es decir se cambió future de 12 horas a 3 horas. En cuanto a las características seleccionadas se mantienen las características previamente seleccionadas para la predicción.

El val_loss del modelo original es: 0.1211 y el nuevo val_loss es de: 0.1224 . Como vemos la diferencia entre ambos valores de pérdida de validación es muy pequeña , o sea que la diferencia es minima, lo cual nos da a entender que tienen un rendimiento muy similar, es normal que haya una pequeña diferencia ya que se ajustó el modelo, sin embargo no se ve afectado de manera significativa.